Strata+ Hadoop WORLD

PRESENTED BY

O'REILLY cloudera

strataconf.com

#StrataHadoop

Building machine-learning apps with Spark

Spark ML and GraphX

Jayant / Vartika / Krishna

Agenda

Overview	5 min (9:00-9:05)	
Lab Environment Setup	10 min (9:05-9:15)	Scala IDE for Eclipse
Spark ML	75 min (9:15-10:15)	Spam Detection, Movie Recommendations, Churn Predictions, Streaming K-Means
Break	30 min (10:30-11:00)	
Spark ML (cont…)	35 min (11:00-11:35)	CrossValidation
GraphX	50 min (11:35-12:25)	Overview, Exploring Structures, Community-Affiliation, Algorithms, The AlphaGo Community, Wikepedia Page Rank
Closing	5 min (12:25-12:30)	

Source Code

- https://github.com/jayantshekhar/strata-2016

Spark ML

Spark ML

Spam Detection	15 min	
Movie Lens Recommendations	15 min	
Streaming K-Means	15 min	
Churn Prediction	15 min	

- Pipeline
  - DataFrame
  - Transformer
  - Estimator
  - Pipeline
- Feature Extractors & Transformers
  - Tokenizer
  - TF/IDF
  - VectorAssembler
  - StringIndexer

- Classification & Regression
  - Logistic Regression
  - Decision Tree
  - Random Forest
  - Gradient Boosted Tree
- Collaborative Filtering
  - ALS
- Frequent Pattern Mining
- Clustering
  - K-Means
  - LDA

Spam Detection

- Logistic Regression

Spam Detection on Enron Dataset

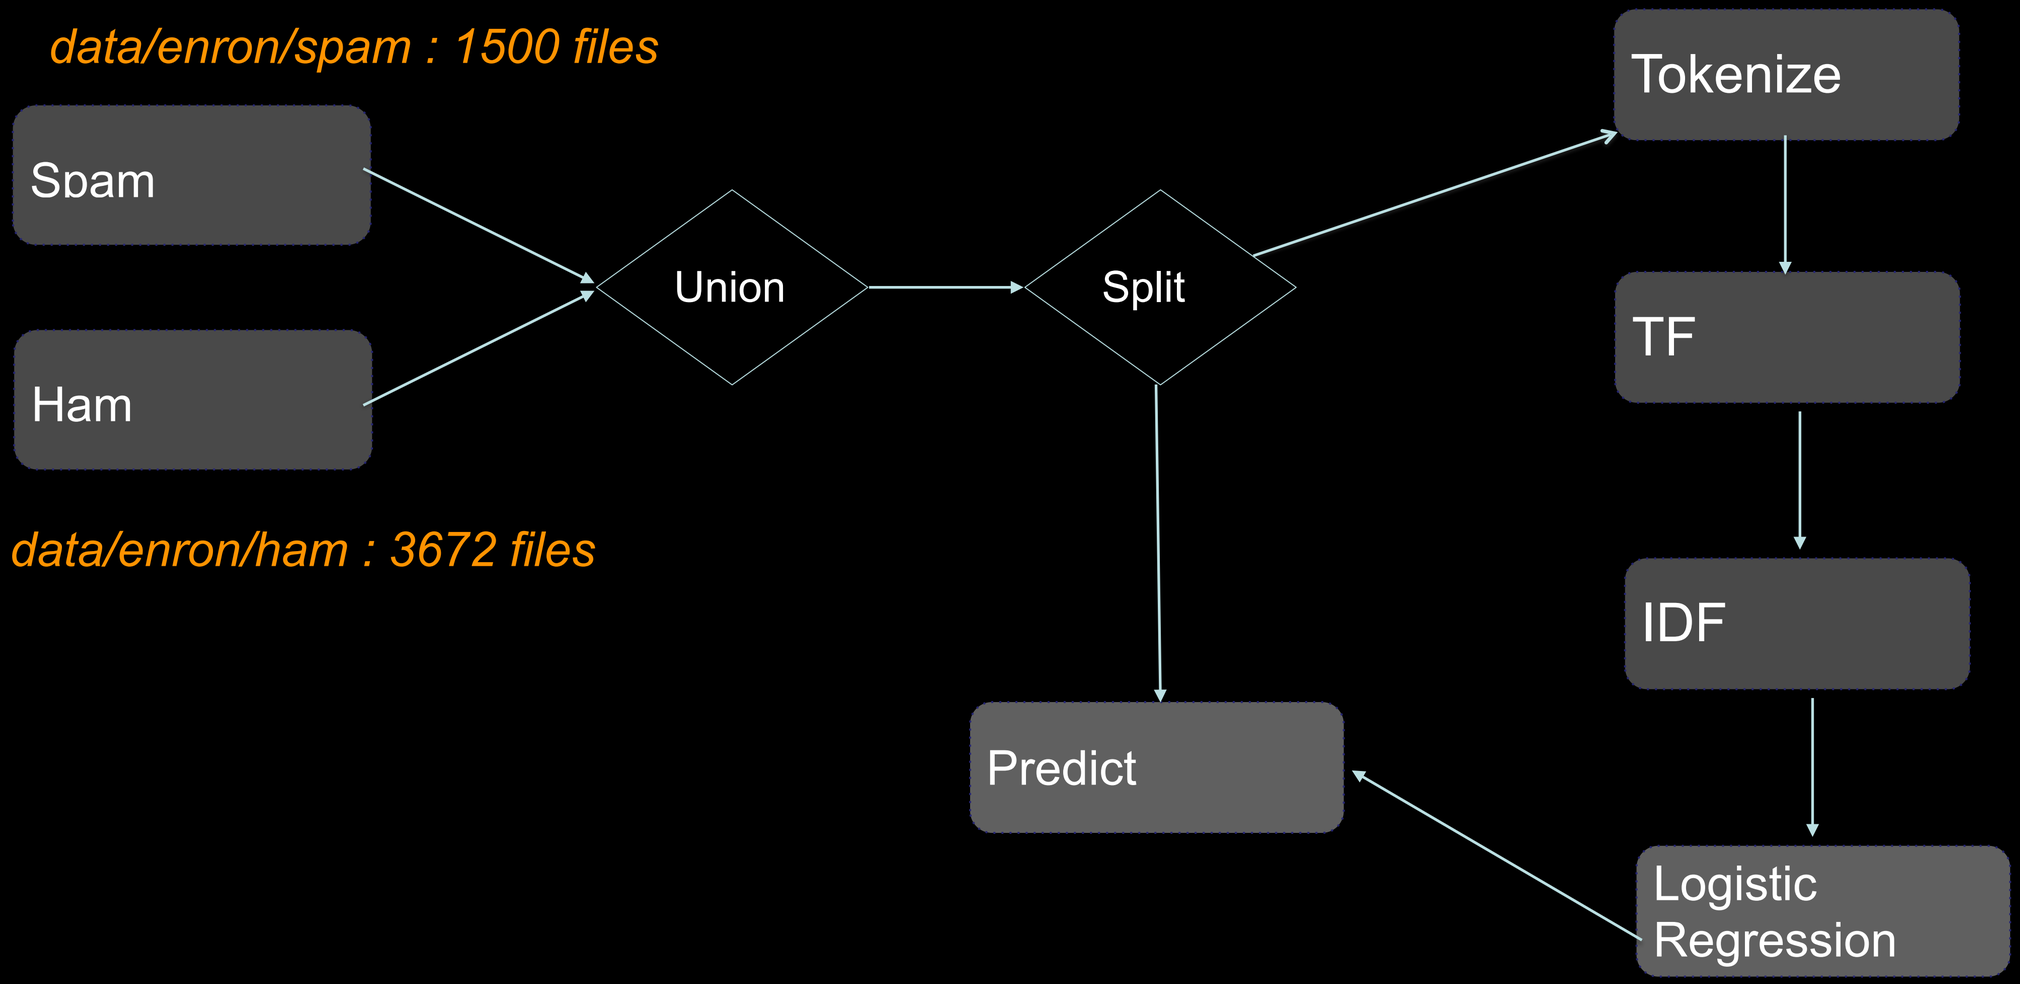

file	text	label
file:/Users/jayan...	Subject: dobmeos ...	1.0
file:/Users/jayan...	Subject: your pre...	1.0
file:/Users/jayan...	Subject: get that...	1.0
file:/Users/jayan...	Subject: await yo...	1.0
file:/Users/jayan...	Subject: coca col...	1.0

file	text	label
file:/Users/jayan...	Subject: christma...	0.0
file:/Users/jayan...	Subject: vastar r...	0.0
file:/Users/jayan...	Subject: calpine ...	0.0
file:/Users/jayan...	Subject: re : iss...	0.0
file:/Users/jayan...	Subject: meter 72...	0.0
file:/Users/jayan...	Subject: mcmullen...	0.0

file	text	label	features	prediction
file:/Users/jayan...	Subject: dobmeos ...	1.0	(262144,[0,33,37,...	1.0
file:/Users/jayan...	Subject: await yo...	1.0	(262144,[0,36,40,...	0.0
file:/Users/jayan...	Subject: real pro...	1.0	(262144,[0,33,36,...	1.0
file:/Users/jayan...	Subject: re : rdd...	1.0	(262144,[0,44,58,...	1.0
file:/Users/jayan...	Subject: cut your...	1.0	(262144,[0,37,39,...	1.0
file:/Users/jayan...	Subject: shut - i...	0.0	(262144,[0,35,40,...	0.0
file:/Users/jayan...	Subject: hpl nomi...	0.0	(262144,[0,38,44,...	0.0
file:/Users/jayan...	Subject: 98 - 673...	0.0	(262144,[0,35,38,...	0.0
file:/Users/jayan...	Subject: hl & p m...	0.0	(262144,[0,33,38,...	0.0
file:/Users/jayan...	Subject: purchasi...	0.0	(262144,[0,34,39,...	0.0
file:/Users/jayan...	Subject: per nels...	0.0	(262144,[0,39,40,...	0.0
file:/Users/jayan...	Subject: see atta...	0.0	(262144,[0,44,58,...	0.0
file:/Users/jayan...	Subject: monthly ...	0.0	(262144,[0,34,36,...	0.0
file:/Users/jayan...	Subject: koch mid...	0.0	(262144,[0,44,46,...	0.0
file:/Users/jayan...	Subject: nom chan...	0.0	(262144,[0,34,40,...	0.0
file:/Users/jayan...	Subject: half day...	0.0	(262144,[0,46,47,...	0.0

Recommendations

- Movie Lens Ratings

MovieLens 100K Dataset

Stable benchmark dataset. 100,000 ratings from 1000 users on 1700 movies. Released 4/1998.

- README.txt
- ml-100k.zip (size: 5 MB, checksum)
- Index of unzipped files

Permalink: http://grouplens.org/datasets/movielens/100k/

MovieLens 1M Dataset

Stable benchmark dataset. 1 million ratings from 6000 users on 4000 movies. Released 2/2003.

- README.txt
- ml-1m.zip (size: 6 MB, checksum)

Permalink: http://grouplens.org/datasets/movielens/1m/

MovieLens 10M Dataset

Stable benchmark dataset. 10 million ratings and 100,000 tag applications applied to 10,000 movies by 72,000 users. Released 1/2009.

- README.html
- ml-10m.zip (size: 63 MB, checksum)

MovieLens

Userid, movie id, rating

0,2,3
0,3,1
0,5,2

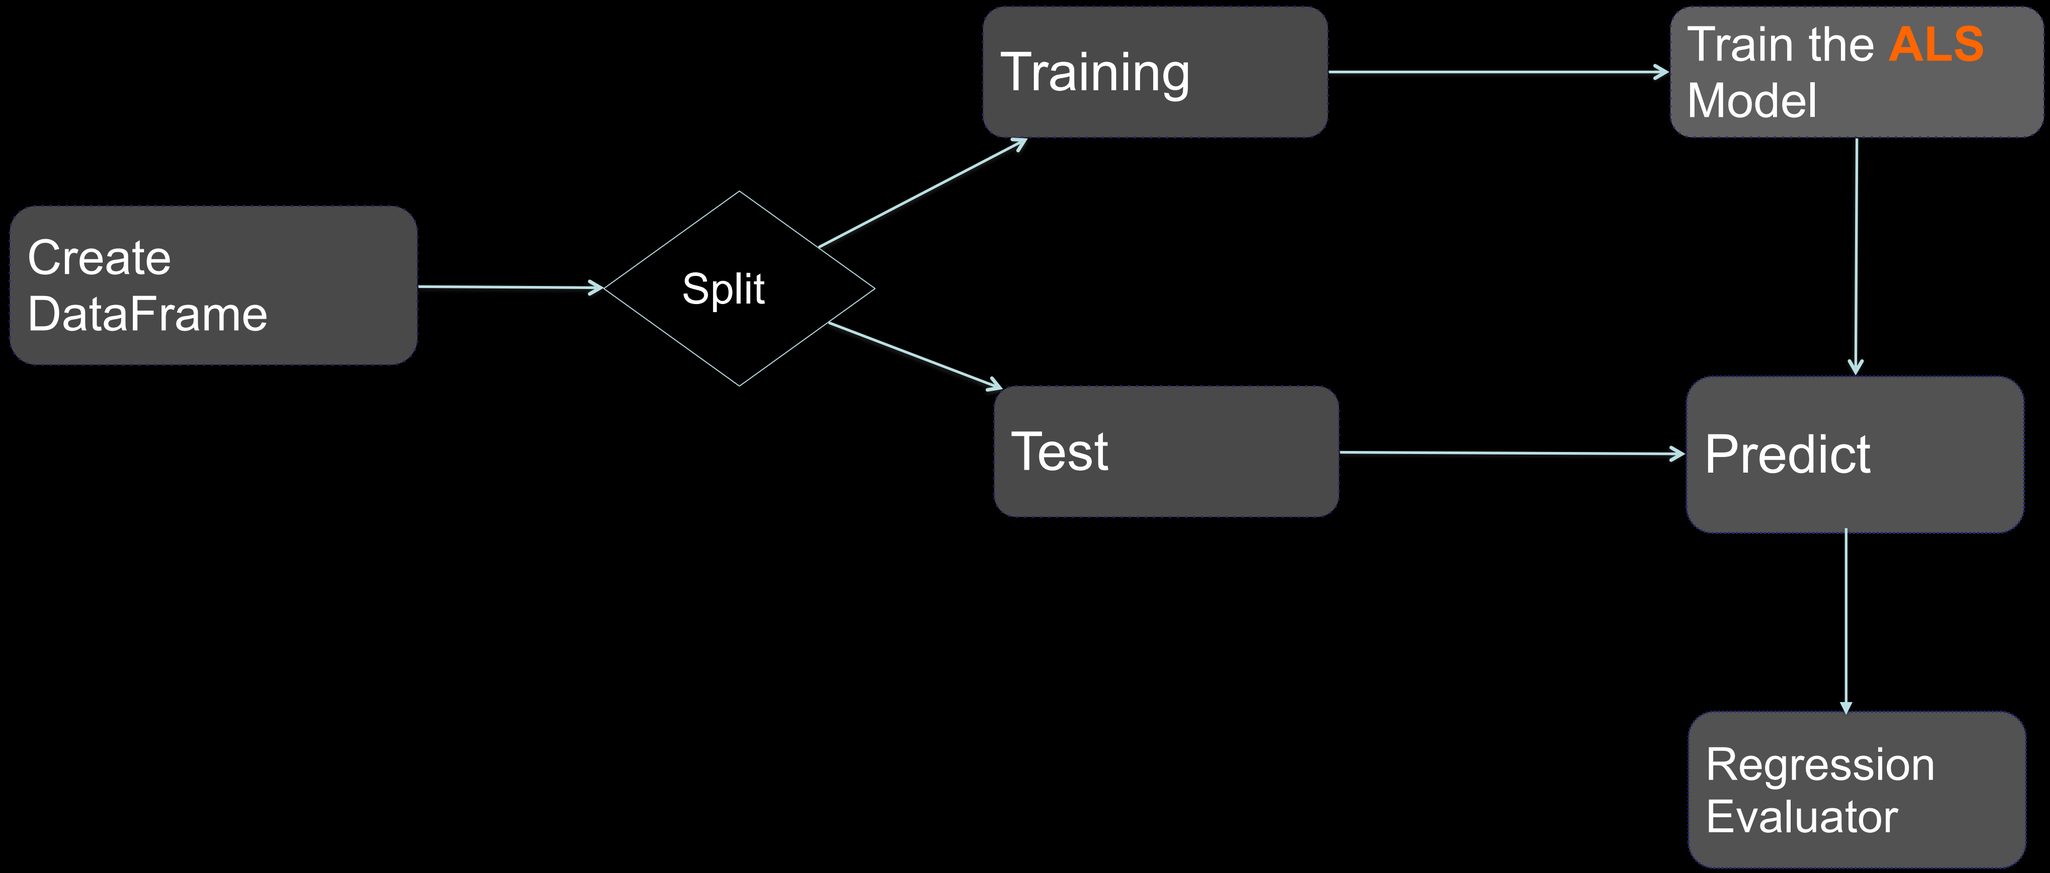

```
root
 |-- C0: string (nullable = true)
 |-- C1: string (nullable = true)
 |-- C2: string (nullable = true)

root
 |-- user: string (nullable = true)
 |-- movie: string (nullable = true)
 |-- rating: string (nullable = true)
```

```
+----+-----+------+
|user|movie|rating|
+----+-----+------+
|   1| 1193|   5.0|
|   1|  661|   3.0|
|   1|  914|   3.0|
|   1| 3408|   4.0|
|   1| 2355|   5.0|
|   1| 1197|   3.0|
|   1| 1287|   5.0|
```

```
|user|movie|rating|prediction|
+----+-----+------+----------+
|5234|   31|   1.0| 2.1774428|
|2242|   31|   5.0| 3.1459289|
|1451|   31|   4.0| 2.3405406|
| 855|   31|   3.0| 2.2023783|
| 855|   31|   3.0| 2.2023783|
|5657|   31|   4.0| 3.7401206|
|5305|   31|   3.0| 2.1689334|
|1306|   31|   3.0| 3.2851512|
```

 * **Alternating Least Squares (ALS) matrix factorization.**
 *
 * ALS attempts to estimate the ratings matrix `R` as the product of two lower-rank matrices,
 * `X` and `Y`, i.e. `X * Yt = R`. Typically these approximations are called 'factor' matrices.
 * The general approach is iterative. During each iteration, one of the factor matrices is held constant, while the other is solved for using least squares. The newly-solved factor matrix is then held constant while solving for the other factor matrix.

Streaming K-Means


Kafka
Flume
HDFS/S3
Kinesis
Twitter
Spark Streaming
HDFS
Databases
Dashboards

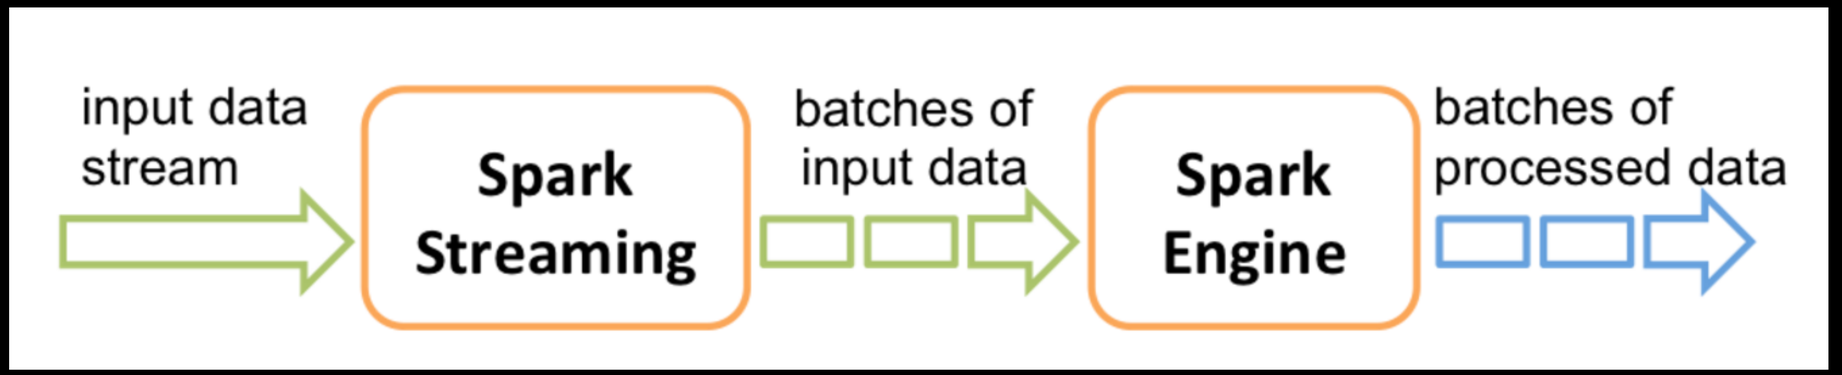
input data stream
Spark Streaming
batches of input data
Spark Engine
batches of processed data

Streaming K-Means

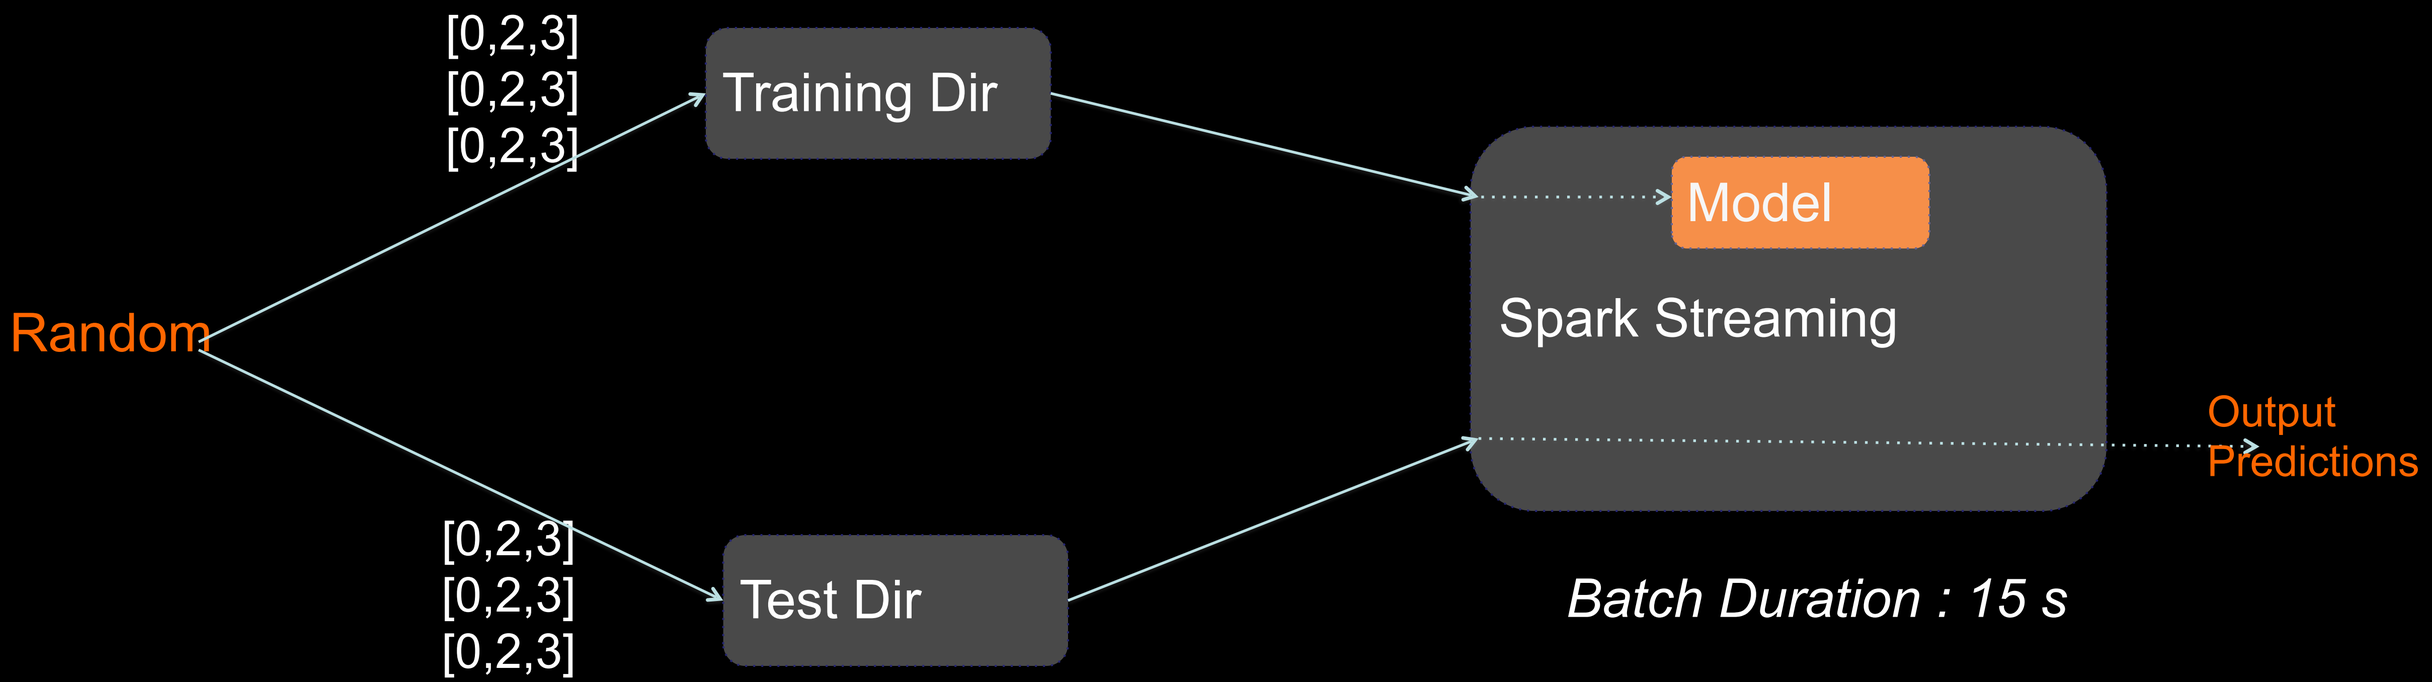

Estimate clusters on one stream of data and make predictions on another stream

Streaming K-Means

- Each point should be formatted as [x1, x2, x3]
- Anytime a text file is placed in ../trainingDir the model would update
- Any time a text file is placed in ../testDir they would be processed to produce predictions using the current model
- The decay can be specified using a halfLife parameter, which determines the correct decay factor such that, for data acquired at time t, its contribution by time t + halfLife will have dropped to 0.5.

```
var trainingDir = "streamingTrainDir"
var testDir = "streamingTestDir"
var batchDuration : Long = 15 // in seconds
var numClusters = 3
```

```
Jayant:strata-2016 jayant$ ls -l streamingDataDir/
total 0
drwxr-xr-x  5 jayant  staff  170 Jun  1 00:09 1
drwxr-xr-x  5 jayant  staff  170 Jun  1 00:10 2
drwxr-xr-x  5 jayant  staff  170 Jun  1 00:10 3
drwxr-xr-x  5 jayant  staff  170 Jun  1 00:10 4
drwxr-xr-x  5 jayant  staff  170 Jun  1 00:10 5
drwxr-xr-x  5 jayant  staff  170 Jun  1 00:11 6
```

```
Jayant:strata-2016 jayant$ ls -l streamingTrainDir/
total 64
-rw-r--r--  1 jayant  staff  6109 Jun  1 00:09 1
-rw-r--r--  1 jayant  staff  6101 Jun  1 00:10 3
-rw-r--r--  1 jayant  staff  6074 Jun  1 00:10 5
-rw-r--r--  1 jayant  staff  6082 Jun  1 00:11 7
```

```
Jayant:strata-2016 jayant$ ls -l streamingTestDir/
total 64
-rw-r--r--  1 jayant  staff  8153 Jun  1 00:10 2
-rw-r--r--  1 jayant  staff  8171 Jun  1 00:10 4
-rw-r--r--  1 jayant  staff  8192 Jun  1 00:11 6
-rw-r--r--  1 jayant  staff  8142 Jun  1 00:11 8
```

Training Data

```
[-0.2887592134448243 6,1.385890499240675 3,0.08997605487060531]
[-1.2805130758440209,0.993958461287273 7,-0.4765545275002676 7]
[-0.010443281100194716,1.439059706483220 7,0.1060992764324971]
[-0.621080758021953,-1.085607452408396 3,-0.6240457792919338]
[0.814710220220870 5,0.3347047775444069,0.998239073229219]
```

Test Data

```
(2.0,[1.314674536226186,-0.5939141316825893,-0.2665244694182238])
(2.0,[0.15896607505959656,-1.3248116154352312,1.7005387494315547])
(1.0,[2.092288692338904,-0.42478085016618355,1.1944557205082678])
(2.0,[0.8231075882687068,-1.7338222010770865,-2.2743871173449 73])
(1.0,[-0.48919972592668304,-1.3353797854575076,1.3845477789028335])
```

Predictions

```
(2.0,1)
(2.0,0)
(1.0,1)
(2.0,1)
(1.0,0)
(2.0,2)
```

- Churn Prediction
- Random Forest

Churn Prediction on Enron Dataset

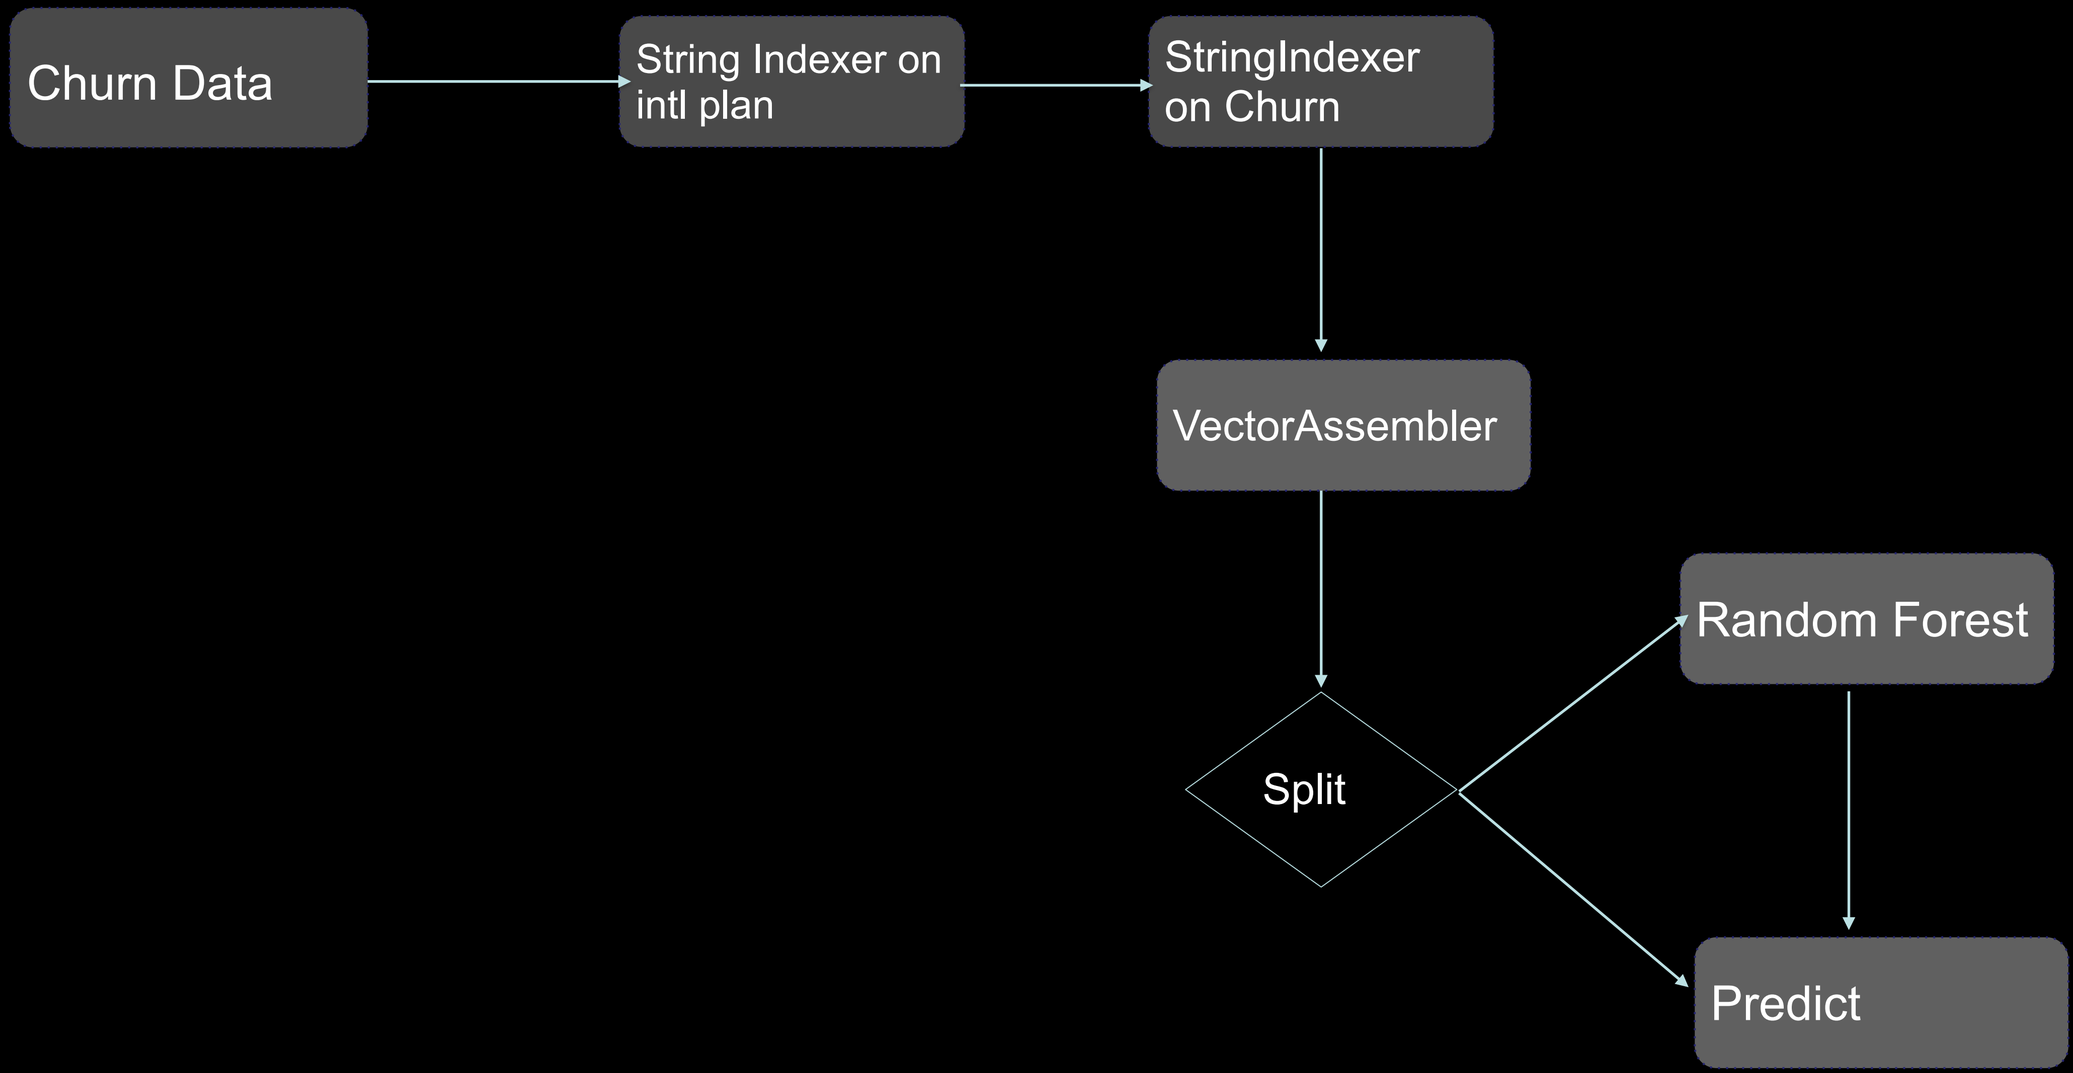

C0	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	C11	C12	C13	C14	C15	C16	C17	C18	C19	C20
KS	128.0	415.0	382-4657	no	yes	25.0	265.1	110.0	45.07	197.4	99.0	16.78	244.7	91.0	11.01	10.0	3.0	2.7	1.0	False.
OH	107.0	415.0	371-7191	no	yes	26.0	161.6	123.0	27.47	195.5	103.0	16.62	254.4	103.0	11.45	13.7	3.0	3.7	1.0	False.
NJ	137.0	415.0	358-1921	no	no	0.0	243.4	114.0	41.38	121.2	110.0	10.3	162.6	104.0	7.32	12.2	5.0	3.29	0.0	False.
OH	84.0	408.0	375-9999	yes	no	0.0	299.4	71.0	50.9	61.9	88.0	5.26	196.9	89.0	8.86	6.6	7.0	1.78	2.0	False.
OK	75.0	415.0	330-6626	yes	no	0.0	166.7	113.0	28.34	148.3	122.0	12.61	186.9	121.0	8.41	10.1	3.0	2.73	3.0	False.

e	account_length	area_code	phone_number	international_plan	voice_mail_plan	number_vmail_messages	total_day_minutes	total_day_calls	total_day_charge	total_eve_minutes	total_eve_calls	total_eve_charge	total_night_mins	total_night_calls	total_night_charge	total_intl_minutes	_intl_calls	total_intl_chargs	num_customer_service_calls	churned
S	128.0	415.0	382-4657	no	yes	25.0	265.1	110.0	45.07	197.4	99.0	16.78	244.7	91.0	11.01	10.0	3	2.7	1.0	False.
H	107.0	415.0	371-7191	no	yes	26.0	161.6	123.0	27.47	195.5	103.0	16.62	254.4	103.0	11.45	13.7	3	3.7	1.0	False.
J	137.0	415.0	358-1921	no	no	0.0	243.4	114.0	41.38	121.2	110.0	10.3	162.6	104.0	7.32	12.2	5	3.29	0.0	False.
H	84.0	408.0	375-9999	yes	no	0.0	299.4	71.0	50.9	61.9	88.0	5.26	196.9	89.0	8.86	6.6	7	1.78	2.0	False.
K	75.0	415.0	330-6626	yes	no	0.0	166.7	113.0	28.34	148.3	122.0	12.61	186.9	121.0	8.41	10.1	3	2.73	3.0	False.
L	118.0	510.0	391-8027	yes	no	0.0	223.4	98.0	37.98	220.6	101.0	18.75	203.9	118.0	9.18	6.3	6	1.7	0.0	False.
A	121.0	510.0	355-9993	no	yes	24.0	218.2	88.0	37.09	348.5	108.0	29.62	212.6	118.0	9.57	7.5	6	2.03	3.0	False.
O	147.0	415.0	329-9001	yes	no	0.0	157.0	79.0	26.69	103.1	94.0	8.76	211.8	96.0	9.53	7.1	6	1.92	0.0	False.

Original Schema…

```
root
 |-- state: string (nullable = true)
 |-- account_length: double (nullable = true)
 |-- area_code: double (nullable = true)
 |-- phone_number: string (nullable = true)
 |-- international_plan: string (nullable = true)
 |-- voice_mail_plan: string (nullable = true)
 |-- number_vmail_messages: double (nullable = true)
 |-- total_day_minutes: double (nullable = true)
 |-- total_day_calls: double (nullable = true)
 |-- total_day_charge: double (nullable = true)
 |-- total_eve_minutes: double (nullable = true)
 |-- total_eve_calls: double (nullable = true)
 |-- total_eve_charge: double (nullable = true)
 |-- total_night_mins: double (nullable = true)
 |-- total_night_calls: double (nullable = true)
 |-- total_night_charge: double (nullable = true)
 |-- total_intl_minutes: double (nullable = true)
 |-- total_intl_calls: double (nullable = true)
 |-- total_intl_chargs: double (nullable = true)
 |-- num_customer_service_calls: double (nullable = true)
 |-- churned: string (nullable = true)
```

Schema after assembler…

```
root
 |-- state: string (nullable = true)
 |-- account_length: double (nullable = true)
 |-- area_code: double (nullable = true)
 |-- phone_number: string (nullable = true)
 |-- international_plan: string (nullable = true)
 |-- voice_mail_plan: string (nullable = true)
 |-- number_vmail_messages: double (nullable = true)
 |-- total_day_minutes: double (nullable = true)
 |-- total_day_calls: double (nullable = true)
 |-- total_day_charge: double (nullable = true)
 |-- total_eve_minutes: double (nullable = true)
 |-- total_eve_calls: double (nullable = true)
 |-- total_eve_charge: double (nullable = true)
 |-- total_night_mins: double (nullable = true)
 |-- total_night_calls: double (nullable = true)
 |-- total_night_charge: double (nullable = true)
 |-- total_intl_minutes: double (nullable = true)
 |-- total_intl_calls: double (nullable = true)
 |-- total_intl_chargs: double (nullable = true)
 |-- num_customer_service_calls: double (nullable = true)
 |-- churned: string (nullable = true)
 |-- label: double (nullable = true)
 |-- international_plan_indx: double (nullable = true)
 |-- features: vector (nullable = true)
```

```
areaUnderROC = 0.6232287449392713
Learned classification forest model:
RandomForestClassificationModel (uid=rfc_3f29d7cd01e1) with 10 trees
  Tree 0 (weight 1.0):
    If (feature 6 <= 133.4)
     If (feature 9 <= 3.0)
      If (feature 0 <= 76.0)
       Predict: 0.0
      Else (feature 0 > 76.0)
       If (feature 5 <= 29.84)
        If (feature 4 <= 64.0)
         Predict: 1.0
        Else (feature 4 > 64.0)
         Predict: 0.0
       Else (feature 5 > 29.84)
        Predict: 1.0
     Else (feature 9 > 3.0)
      If (feature 2 <= 0.0)
       Predict: 0.0
      Else (feature 2 > 0.0)
       Predict: 1.0
    Else (feature 6 > 133.4)
     If (feature 3 <= 272.6)
      If (feature 5 <= 12.67)
       If (feature 0 <= 64.0)
        Predict: 0.0
       Else (feature 0 > 64.0)
        If (feature 6 <= 180.6)
         Predict: 1.0
        Else (feature 6 > 180.6)
```

prediction	label	features
0.0	0.0	[93.0,0.0,0.0,271...
0.0	0.0	[95.0,0.0,0.0,238...
0.0	0.0	[75.0,0.0,0.0,166...
0.0	0.0	[116.0,0.0,34.0,2...
1.0	1.0	[151.0,1.0,0.0,21...
0.0	0.0	[68.0,0.0,0.0,237...
0.0	0.0	[107.0,0.0,0.0,13...
0.0	0.0	[141.0,0.0,32.0,1...
0.0	1.0	[159.0,1.0,0.0,25...

discover actual shopping behavior

- Frequent Pattern Mining
- FPG

Frequent Pattern Mining

- Mllib has parallel implementation of FP-Growth

  - minSupport: the minimum support for an itemset to be identified as frequent. For example, if an item appears 3 out of 5 transactions, it has a support of 3/5=0.6.
  - numPartitions: the number of partitions used to distribute the work.

FPGrowth

Create RDD of ArrayList<String>

r z h k p
z y x w v
u t s
s x o n r
x z y m t
s q e

Run FPGrowth

[s], 3
[s,x], 3
[s,x,z], 2
[s,z], 2
[r], 3
[r,x], 2
[r,z], 2
[y], 3
[y,s], 2
[y,s,x], 2

Print Results

ML Pipelines

Spark ML

- DataFrames
- Transformer
- Estimator
- Pipeline

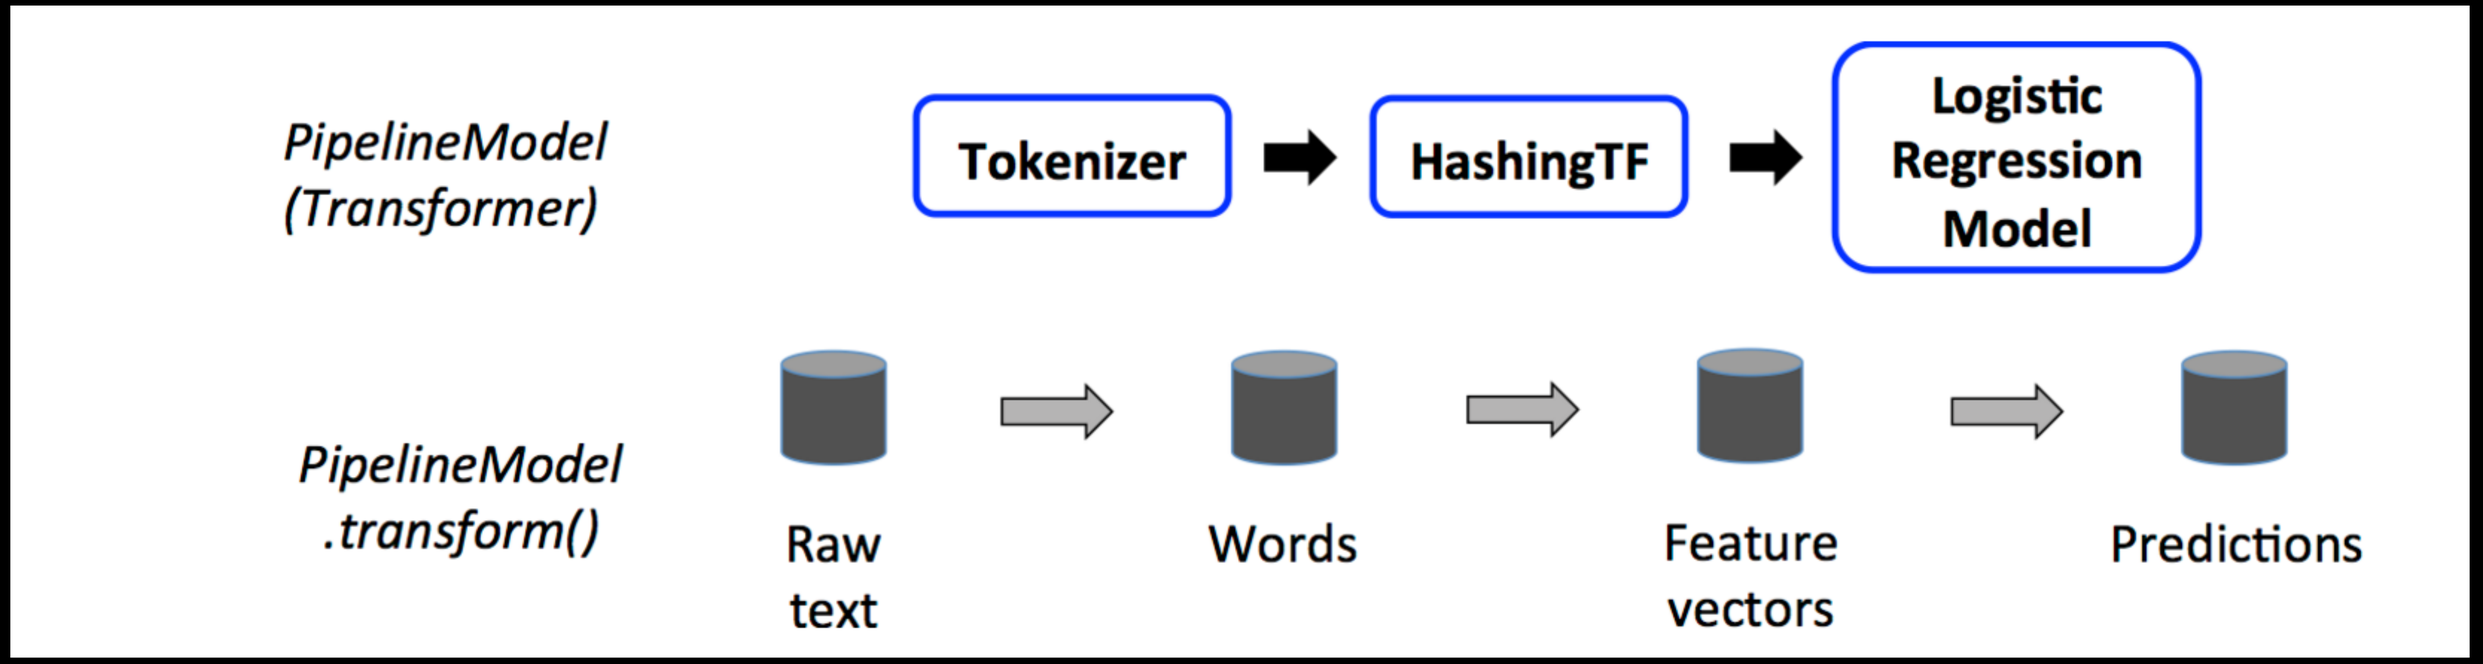

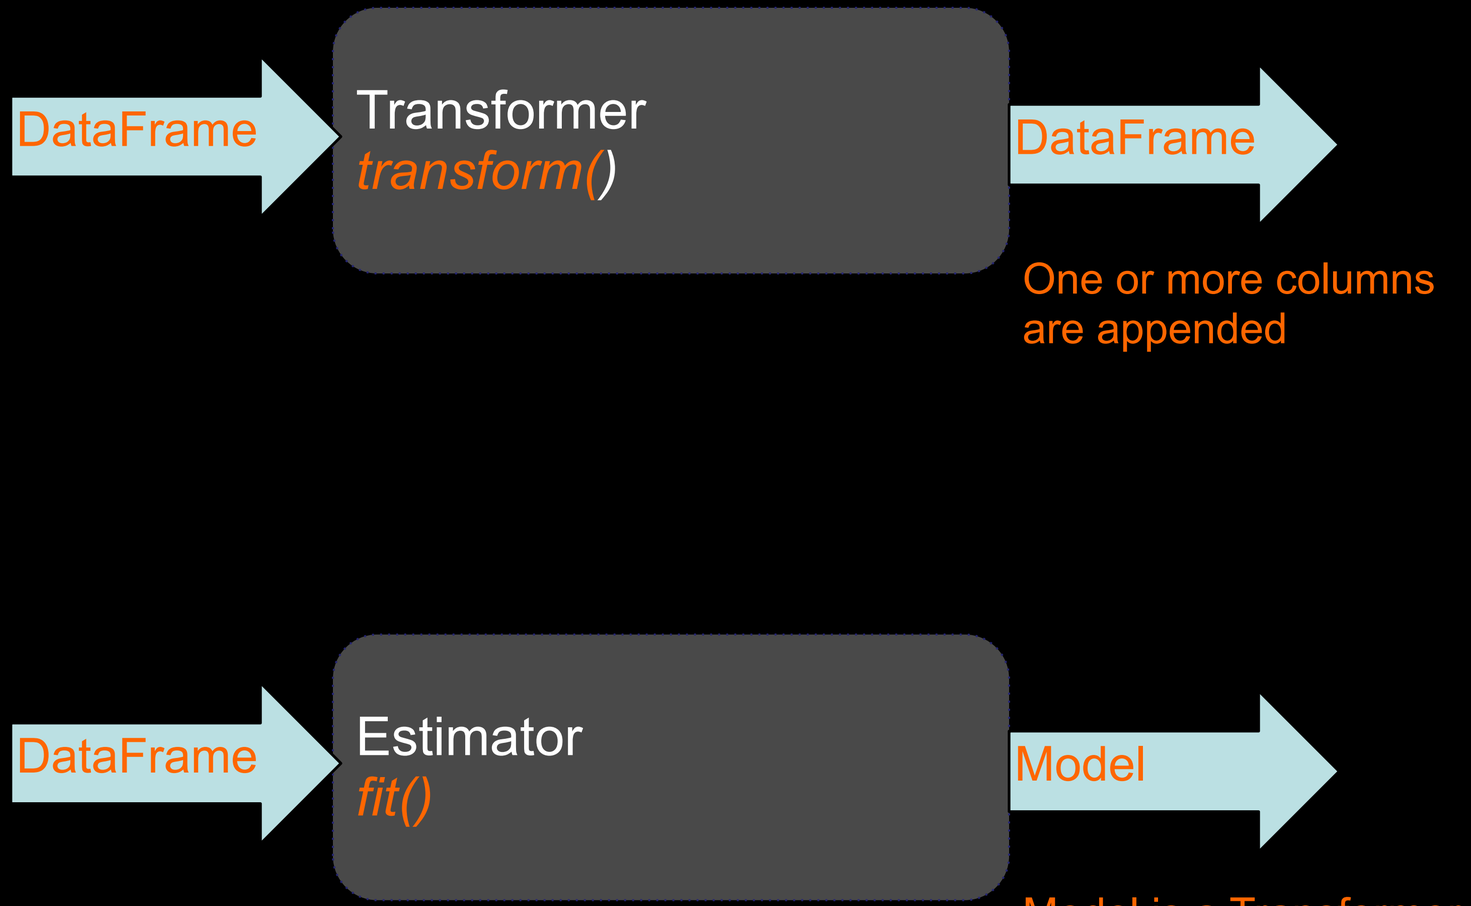
DataFrame
Transformer
transform()
DataFrame
One or more columns are appended
DataFrame
Estimator
fit()
Model
Model is a Transformer

- Titanic Survival Prediction
- Random Forest

Titanic

Remove the header row when reading

- **Data**

PassengerId,Survived,Pclass,Name,Sex,Age,SibSp,Parch,Ticket,Fare,Cabin,Embarked

1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S

2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC 17599,71.2833,C85,C

3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S

- Target Variable
  - Survived
- Predictor Variables
  - Pclass, Sex, Age, Fare

Titanic DataSet

VARIABLE DESCRIPTIONS:

survival	Survival
	(0 = No; 1 = Yes)
pclass	Passenger Class
	(1 = 1st; 2 = 2nd; 3 = 3rd)
name	Name
sex	Sex
age	Age
sibsp	Number of Siblings/Spouses Aboard
parch	Number of Parents/Children Aboard
ticket	Ticket Number
fare	Passenger Fare
cabin	Cabin
embarked	Port of Embarkation

(C = Cherbourg; Q = Queenstown; S = Southampton)

NOTES:

Pclass is a proxy for socio-economic status (SES)
1st ~ Upper; 2nd ~ Middle; 3rd ~ Lower

Age is in Years; Fractional if Age less than One (1)
If the Age is Estimated, it is in the form xx.5

Titanic

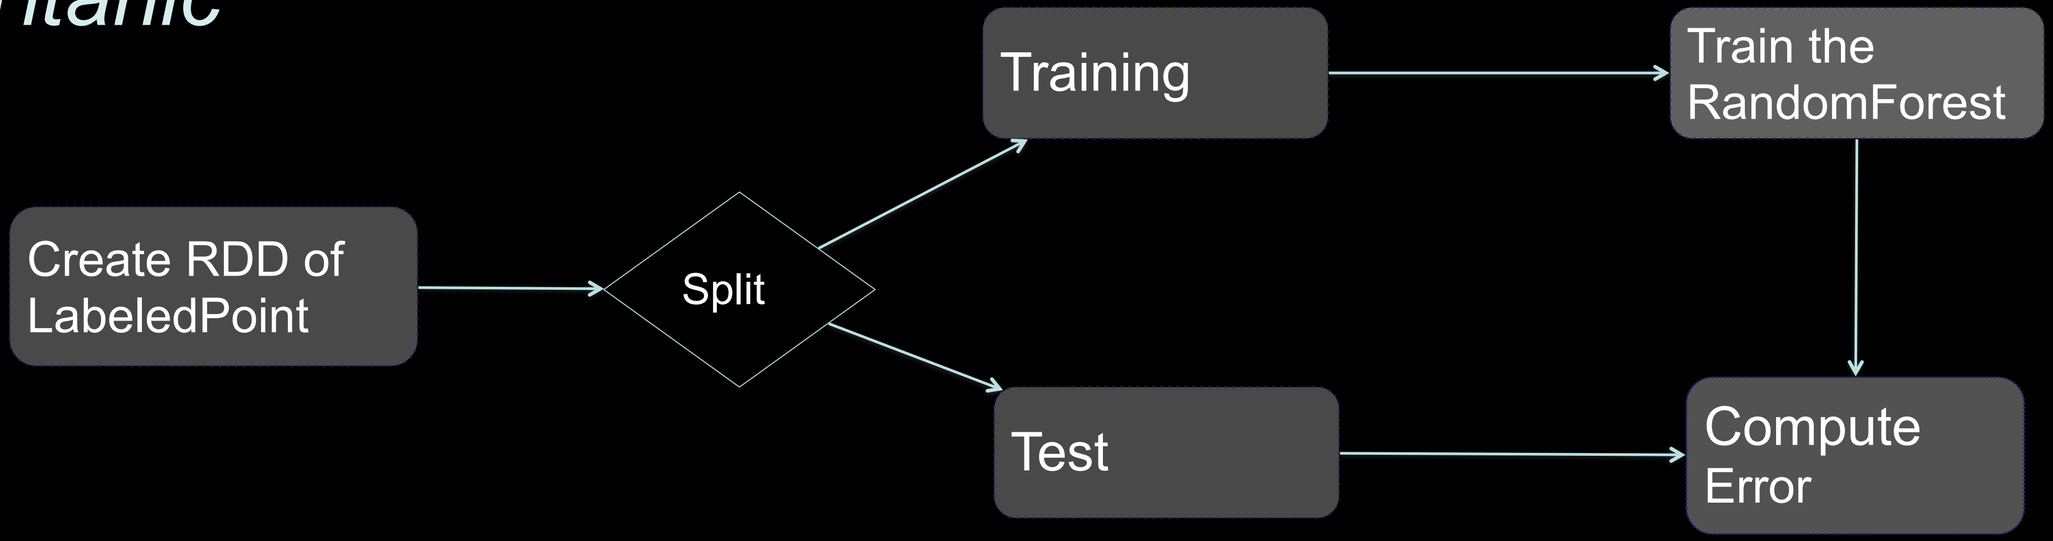


```
root
 |-- PassengerId: string (nullable = true)
 |-- Survived: string (nullable = true)
 |-- Pclass: string (nullable = true)
 |-- Name: string (nullable = true)
 |-- Sex: string (nullable = true)
 |-- Age: string (nullable = true)
 |-- SibSp: string (nullable = true)
 |-- Parch: string (nullable = true)
 |-- Ticket: string (nullable = true)
 |-- Fare: string (nullable = true)
 |-- Cabin: string (nullable = true)
 |-- Embarked: string (nullable = true)
```

Random Forest

- **numTrees**: Number of trees in the forest.
- **maxDepth**: Maximum depth of each tree in the forest.
- **categoricalFeaturesInfo**: Specifies which features are categorical and how many categorical values each of those features can take. This is given as a map from feature indices to feature arity (number of categories). Any features not in this map are treated as continuous.
  - E.g., Map(0 -> 2, 4 -> 10) specifies that feature 0 is binary (taking values 0 or 1) and that feature 4 has 10 categories (values {0, 1, ..., 9}). Feature indices are 0-based: features 0 and 4 are the 1st and 5th elements of an instance's feature vector.

- Tree 0:
- If (feature 0 in {0.0})
- If (feature 4 <= 8.7125)
- If (feature 3 <= 0.0)
- If (feature 2 <= 0.0)
- Predict: 0.0
- Else (feature 2 > 0.0)
- Predict: 0.0
- Else (feature 3 > 0.0)
- If (feature 1 <= 0.42)
- Predict: 1.0
- Else (feature 1 > 0.42)
- Predict: 0.0
- Else (feature 4 > 8.7125)
- If (feature 1 <= 14.0)
- If (feature 2 <= 2.0)
- Predict: 1.0
- Else (feature 2 > 2.0)
- Predict: 0.0
- Else (feature 1 > 14.0)

- Tree 1:
- If (feature 0 in {0.0})
- If (feature 4 <= 9.8375)
- If (feature 4 <= 7.8958)
- If (feature 4 <= 7.8292)
- Predict: 0.0
- Else (feature 4 > 7.8292)
- Predict: 0.0
- Else (feature 4 > 7.8958)
- If (feature 2 <= 0.0)
- Predict: 0.0
- Else (feature 2 > 0.0)
- Predict: 1.0
- Else (feature 4 > 9.8375)
- If (feature 3 <= 0.0)
- If (feature 4 <= 26.0)
- Predict: 0.0
- Else (feature 4 > 26.0)
- Predict: 0.0
- Else (feature 3 > 0.0)